MACHINE LEARNING (SPRING 2025)

Name : Bandari Madhuri

UID: U01093236

Email : bandari.24@wright.edu

This **dataset** includes follow-up data on cases of breast cancer, with an emphasis on whether patients had a cancer recurrence or stayed disease-free. The Outcome variable specifies whether the cancer returned ("R") or did not ("N"), and each instance is uniquely identified by its ID. For patients who **have experienced a recurrence (R)**, the Time variable indicates the time until recurrence; for those who have n**ot experienced a recurrence (N)**, it represents the time until disease-free survival.

Radius, Texture, Area, Concavity, Concave Points, Smoothness, Compactness, Symmetry, and Fractal Dimension are some of these characteristics. Three statistical measures—the mean (average value across all nuclei), standard error (variation in values), and worst (mean of the three largest values)—were calculated for each of these attributes.

Furthermore, two significant clinical characteristics are noted: **Lymph Node Status**, which indicates the quantity of impacted lymph nodes at the time of surgery, and Tumor Size, which quantifies the tumor's dimension in centimeters. These characteristics make this dataset useful for medical research and predictive modeling since they offer crucial insights into the severity and course of breast cancer.

Exploratory Data Analysis (EDA)

Exploratory Data Analysis (EDA) is an important phase in machine learning since it helps us comprehend data distributions, uncover trends, and identify potential abnormalities. In this study, we examine the summary statistics of important numeric variables from the given breast cancer dataset.

A) Summarize the statistics of these variables:

1. **Count**: Each variable has **198 observations**, indicating that all features and few features (mean_radius, mean_texture, lymph_node_status) have data with missing values.
2. Central Tendency (Mean & Median): The **mean** values represent the average measurement for each feature and the **median (50% percentile)** gives the middle value, which helps in understanding the distribution.
3. Variability (Standard Deviation, Min, Max, Percentiles): **Standard deviation (std)** measures how spread out the values are, **Min and max** show the range of values, **25%, 50%, 75% percentiles** help understand how the data is distributed.

4. **Numeric variables in the dataset:** mean_radius, mean_texture, mean_perimeter, mean_area, mean_smoothness, mean_compactness, mean_concavity and mean_concave_points.

The summary statistics of the numeric variables provide key insights into the dataset's distribution and variability: The count values indicate that some features, such as **mean_radius** and **mean_texture**, have slightly fewer observations (194) than others (198), suggesting possible missing data. The **mean and median values** for most features are close, indicating a relatively **symmetrical distribution**, though some features exhibit skewness. Features like **mean_area (970.04) and mean_perimeter (114.85)** show high variability, as reflected in their large standard deviations (352.15 and 21.38, respectively), indicating a wide range of tumor sizes. Additionally, certain variables, such as **mean_area (max = 2250) and mean_concavity (max = 0.4268)**, have extreme maximum values that significantly exceed their **75th percentile values**, suggesting the presence of potential outliers. The shape-based features, including **mean_compactness, mean_concavity, and mean_concave points**, exhibit moderate variability but lower mean values compared to size-related features. Given these observations, further exploration of feature correlations and visualizations such as histograms and box plots would be beneficial in identifying relationships and patterns within the dataset.

B) **Categorical Variable "Outcome":** The **Outcome** variable in this dataset indicates whether a patient experienced a **recurrence of cancer ("R")** or remained **disease-free ("N")** during the follow-up period. A statistical summary of this variable provides key insights into its distribution.

Count: The total number of cases recorded in the dataset.
Unique Values: The number of distinct categories in the outcome variable, which are expected to be **two ("R" and "N")**.
Top Value: The most frequently occurring category, which helps identify the dominant class in the dataset.
Frequency of Top Value: The number of occurrences of the most common category, which is useful in understanding whether the dataset is balanced or imbalanced.

This summary is crucial for evaluating **class distribution**, as an imbalanced dataset—where one category significantly outweighs the other—can affect model performance. If necessary, techniques like **resampling or weighted learning** may be applied to address imbalance before training a machine learning model. Additionally, since the outcome variable is categorical, it needs to be converted into a numerical format (e.g., **"N" → 0 and "R" → 1**) for use in machine learning models.

The dataset's outcome variable indicates if a patient experienced a cancer recurrence ("R") or stayed cancer-free ("N"). The total number of cases is **198**, and there are no missing values. "N" was the most frequent result among these, as most patients **(151 instances) did not have a recurrence**. **Unique Values**: There are **two distinct categories** in the outcome variable (**"N" and "R"**). **Most Frequent Category (Top Value)**: The most common outcome is **"N"**, meaning a larger number of patients did not experience cancer recurrence. **Frequency of Top Value**: The

"N" category appears **151 times**, indicating that many cases in the dataset belong to the **no recurrence** group.

C) The **dataset's Outcome variable** is a categorical variable that indicates whether a patient had a cancer recurrence ("R") or cancer-free ("N"). This variable must be transformed into a numerical format since ML_algorithms need numerical data to process. Assigning "**R" = 1 (Recurrence)** and **"N" = 0 (No Recurrence)** is an easy and efficient method of encoding it. Following encoding, a new column called "outcome_encoded" is added to the dataset, in which every **"N" is changed to 0 and every "R" to 1**. This change maintains the variable's original meaning while enabling its use in predictive modeling. The dataset is prepared for additional analysis and model training once the first few rows verify that the encoding was applied appropriately.

D) The dataset contains several features that are highly correlated, meaning they provide very similar information. For example, **mean_radius and mean_perimeter** have a correlation of **0.9959**, and **mean_area and mean_radius** have a correlation of **0.9929**. This strong relationship suggests that one of these features can be removed without losing significant information. Based on the experiments so far, it is evident that features like **mean_radius, mean_perimeter, and mean_area** contain overlapping information. Similarly, **worst_radius, worst_area, and worst_perimeter** show high correlations, indicating redundancy. Removing one feature from each of these highly correlated pairs will not impact on the overall analysis but will help in creating a more efficient and interpretable model.

E) The correlation between **mean_perimeter** and **se_perimeter** in the dataset is **0.6099**, indicating a **moderate positive correlation**. This suggests that as the mean perimeter of tumors increases, the standard error (SE) of the perimeter also tends to increase.

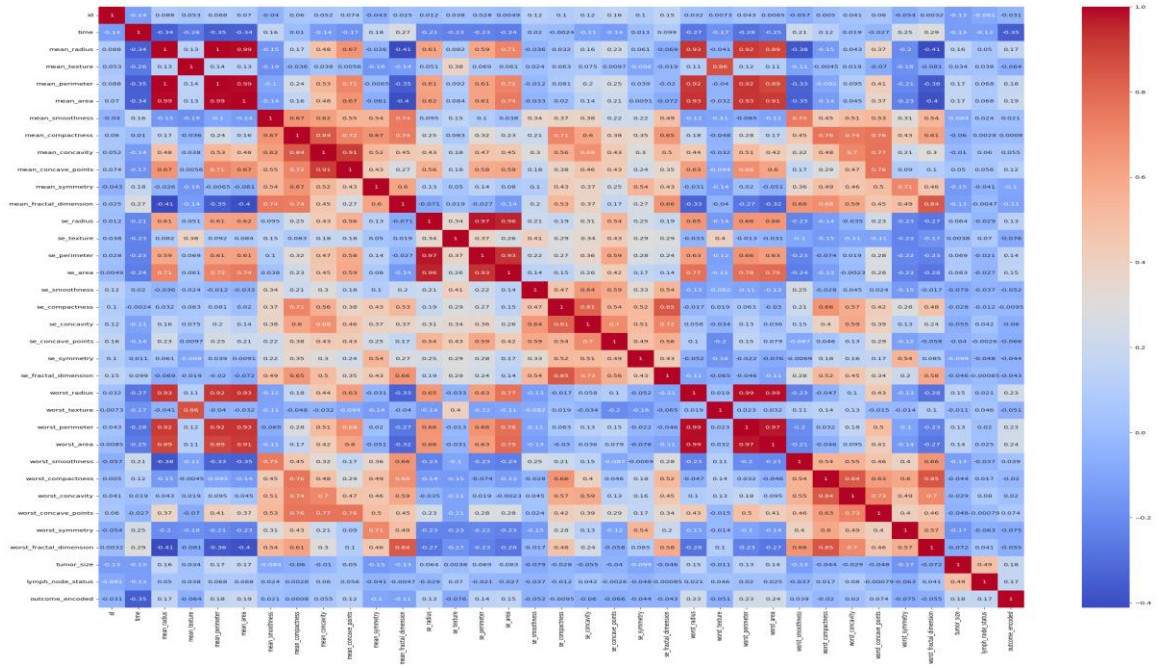

The **above correlation heatmap visually** represents feature relationships in the dataset, with red indicating **positive correlations** and blue **indicating negative ones**. The diagonal shows perfect self-correlations (1). There appears to be a moderately positive connection between **mean_perimeter and se_perimeter (0.6099),** indicating that higher standard errors are linked to larger tumor perimeters. This aids in the selection of features and enhances model performance.

Logistic Regression with One Variables

1. Linear regression with **mean_area** as the predictor provides an initial approach to understanding **breast cancer recurrence**. While it helps capture the trend, its effectiveness depends on the strength of correlation between **mean_area** and recurrence.
2. The **confusion matrix** aids in performance evaluation, highlighting areas where the model excels and where improvements are needed. **True Positives (TP):** Correctly predicted recurrence cases, **True Negatives (TN):** Correctly predicted non-recurrence cases, **False Positives (FP):** Incorrectly predicted recurrence cases, **False Negatives (FN):** Missed recurrence cases.
3. From the confusion matrix, we compute essential metrics such as **accuracy, precision, recall, and F1-score**, which help evaluate the model's effectiveness.

Explanations Graphs

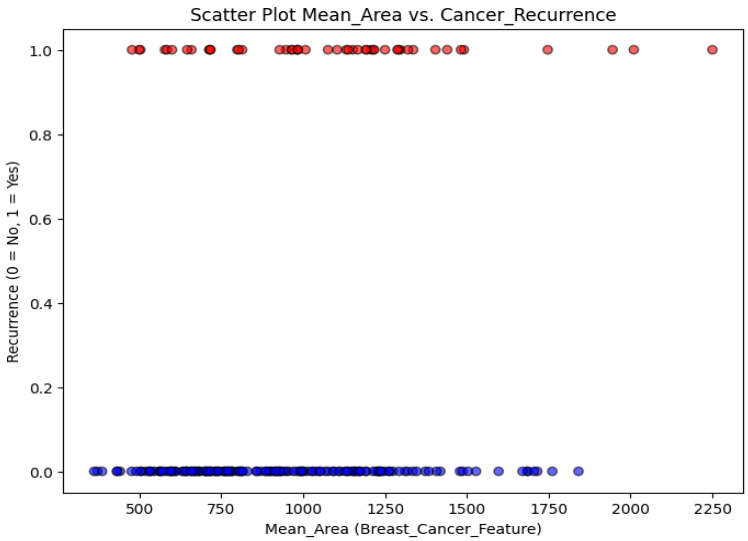

The scatter plot of **mean_area** vs. **cancer recurrence** suggests a trend where higher **mean_area** values are associated with recurrence (indicated by points at 1). However, there is a significant overlap, meaning a single feature may not be sufficient for accurate classification. The **blue points (0)** indicate **non-recurrence**, while the **red points (1)** indicate recurrence. While there is some clustering of recurrence cases at higher mean_area values, there is significant **overlap between recurrence and non-recurrence cases**, suggesting that using only this feature may not provide a clear decision boundary.

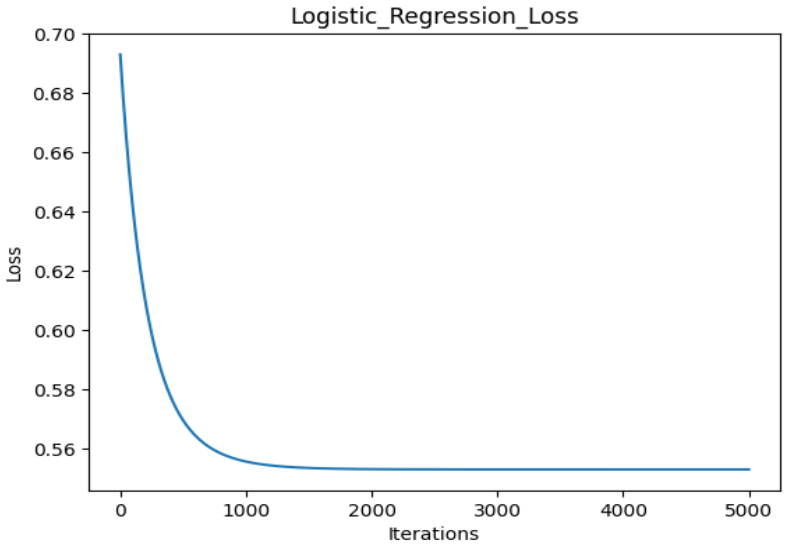

The logistic regression loss **curve shows a steady decrease in loss**, converging to a stable value, indicating successful optimization of the model parameters. The **logistic regression loss curve** demonstrates how the **model optimizes its parameters over multiple iterations**. Initially, the **loss is high**, but it **steadily decreases**, indicating that the model is learning. The curve flattens at a lower value, signifying convergence and stable model parameters.

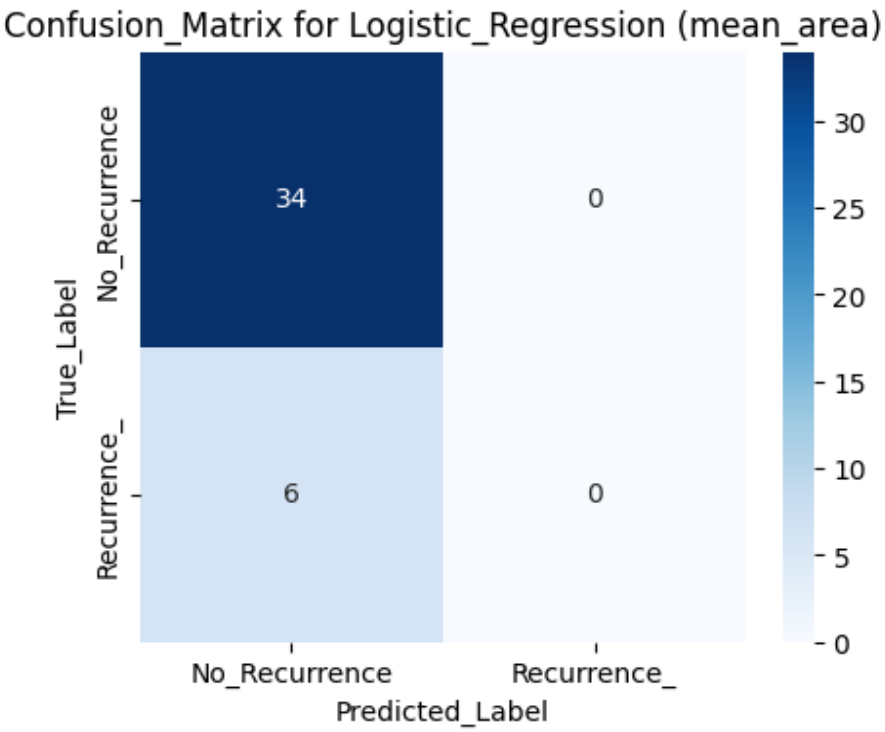

```
Confusion Matrix:
                Predicted No   Predicted Yes
Actual No     |    34         0  |  (TN  FP)
Actual Yes    |    6          0  |  (FN  TP)
```

```
Final Model Parameters (Theta__): [-1.09287264  0.43620237]
Accuracy using one varaible: 85.00%
Precision using one varaible: 0.00%
Recall using one varaible: 0.00%
F1 Score using one varaible: 0.00%
```

According to the results of the **confusion matrix**, the model correctly identified 34 non-recurrence cases (True Negatives, TN) but was unable to accurately predict any recurrence cases (False Negatives, FN). Although the **accuracy was high at 85%,** the model's precision, recall, and F1-score were all 0.00%, indicating that it is biased toward the majority class and cannot anticipate recurrent cases. The model achieved **85% accuracy**, but this is misleading because it only correctly predicts the majority class (non-recurrence).The recall of **0%** means that the model is completely ineffective in identifying recurrence cases. The **model performed poorly** in detecting recurrent cases, as evidenced by its zero precision and recall, despite its high accuracy. The findings imply that relying solely on "**mean_area**" as a predictor is inadequate, and **adding more characteristics could enhance the model's** functionality. **Using only "mean_area" is not sufficient** for predicting recurrence. **More features should be included** in the model to improve performance.

Logistic Regression with Multiple Variables

1. The logistic regression model uses **12 specific variable**s to predict the recurrence of breast cancer. The dataset is preprocessed using **normalization** and **missing value** imputation.
2. Gradient Descent ensures loss minimization across iterations by optimizing the model parameters.
3. The trained model's **accuracy: precision, recall, F1-score, and a confusion matrix** are used to assess performance. With visualization supporting analysis, the results show how well logistic regression works for binary classification in medical diagnosis.
4. **12_Variables:** mean_radius, mean_texture, mean_perimeter, mean_area, mean_smoothness, mean_compactness, mean_concavity, mean_concave_points, mean_fractal_dimension, se_perimeter, se_texture, se_area.

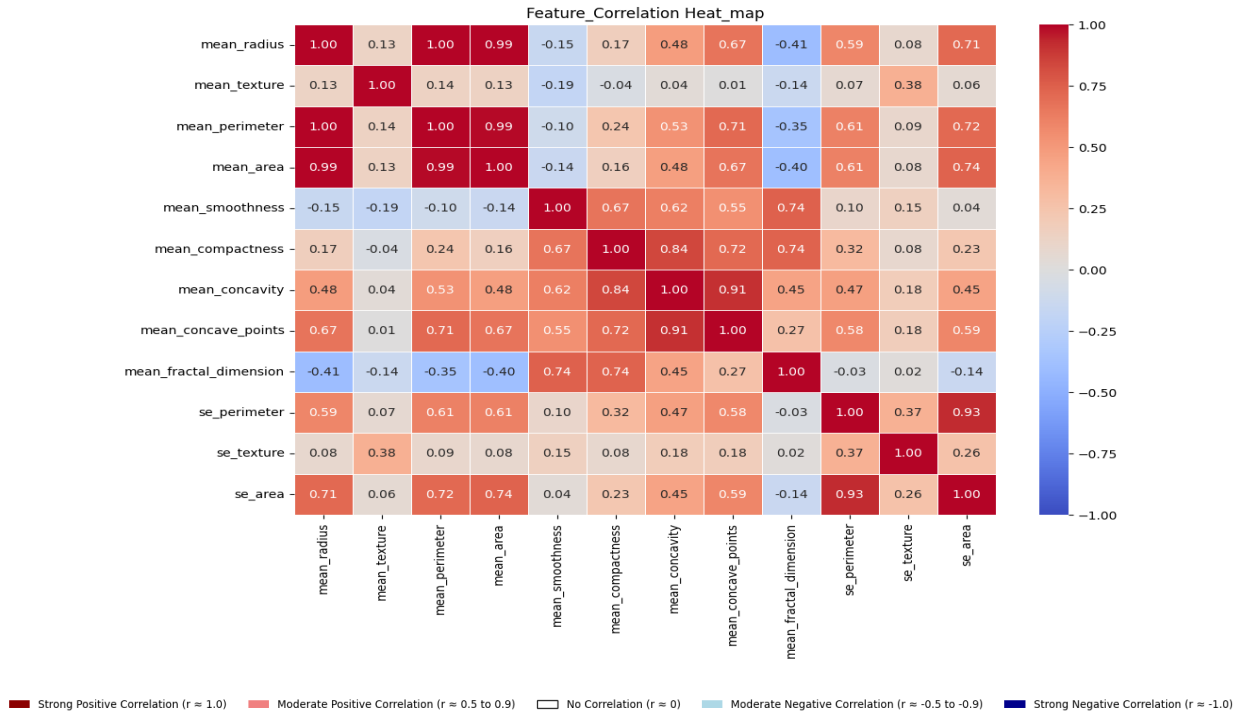

The **heatmap graph visualizes** the correlation matrix of the **12 selected features**. Each cell represents the correlation coefficient between a pair of features, with colors indicating both the strength and direction of their relationship. **Dark red** hues suggest a **strong positive correlation**, while **dark blue** indicates a **strong negative correlation; orange represents** moderate Positive_correlation; colors closer to **white** imply little to **no correlation and light blue** represents moderate negative correlation. This visual tool helps quickly identify **highly correlated features**, which could imply redundancy issues that might impact the performance and interpretability of the **logistic regression model.**

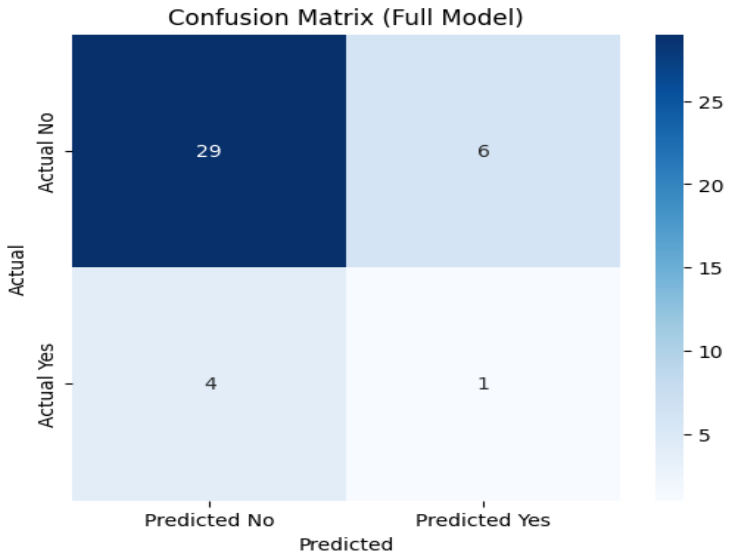

This graph shows the distribution of expected and actual results. Out of all the predictions, 29 cases were accurately classified as negatives (true negatives), according to the matrix, whereas 6 negatives were mistakenly labeled as positives (false positives). On the other hand, just one positive (true positive) was accurately detected, while four positives were mistakenly categorized as negatives (false negatives).

```
Confusion Matrix (Full Model):
                Predicted No   Predicted Yes
Actual No     |    29         6  |  (TN  FP)
Actual Yes    |    4          1  |  (FN  TP)
Accuracy: 75.00%
Precision: 14.29%
Recall: 20.00%
F1 Score: 16.67%
```

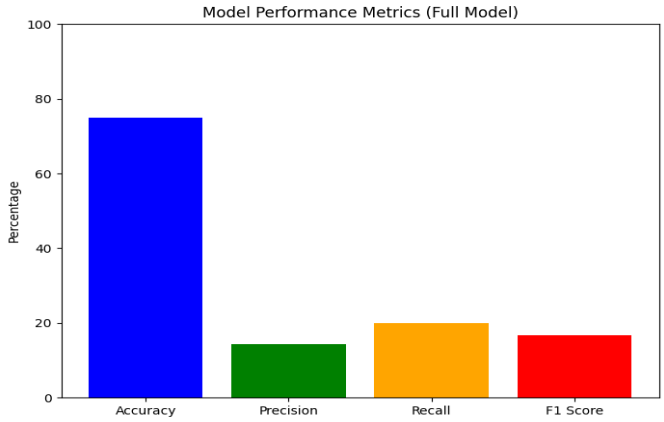

The above graph model performance metric summarizes evaluation metrics—**accuracy, precision, recall, and F1 score**—in a bar format. While the model achieves a moderate accuracy of 75%, the precision (14.29%), recall (20%), and F1 score (16.67%) are significantly lower. This discrepancy highlights that the model, despite **correctly classifying many negatives**, struggles considerably with the positive class, indicating issues such as class imbalance or inadequate feature representation.

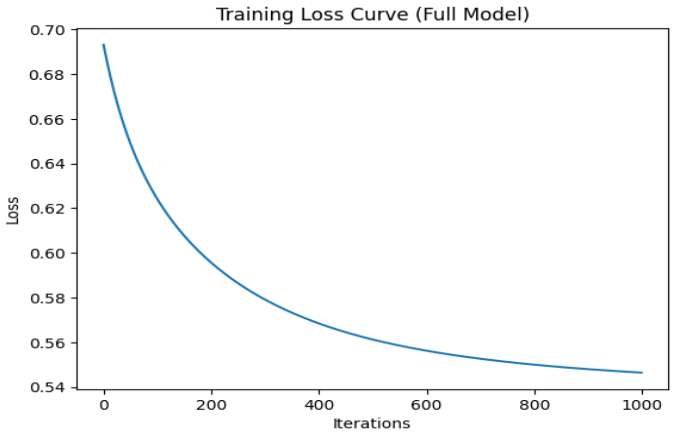


```
Final Model Parameters (Theta for 12 Features): [-0.87826473  0.05102773  0.02977896  0.0892346   0.13455434  0.16660577
 -0.09965706 -0.03910959  0.05582143 -0.1752941   0.2095278  -0.0509459
  0.08732987]
```

The **training loss curve** illustrates how the cost function, or loss, drops when the gradient descent algorithm is run through more iterations. This graph's downward trend indicates that the model is learning, and that the optimization process is ending. Even though the loss diminishes with time, the poor performance metrics suggest that the model may be underfitting the positive class or failing to adequately capture the underlying patterns.

Logistic Regression using Forward Selection Variables

1. Breast cancer recurrence prediction is crucial for improving patient outcomes.
2. The study uses **Logistic Regression** with **forward selection** to identify the most significant features for predicting recurrence.
3. The dataset includes **12 features**, such as mean radius, texture, area, and perimeter.
4. Forward selection helps in choosing the best subset of features that enhance model performance.
5. The effectiveness of the model is evaluated using **accuracy, precision, recall, and F1-score.**
6. The selected features are analyzed to understand their role in recurrence prediction, contributing to better diagnostic tools in medical research.

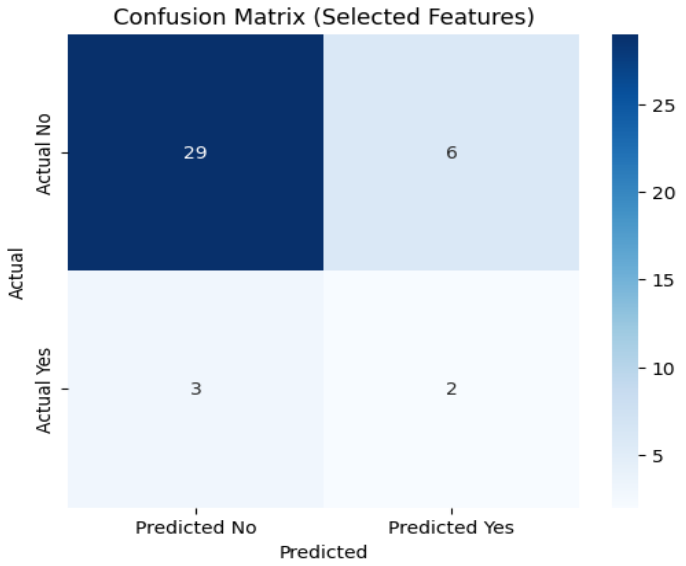


```
Selected Features: ['mean_area', 'se_perimeter', 'mean_fractal_dimension', 'mean_smoothness', 'mean_compactness', 'se_texture', 'mean_texture']

Confusion Matrix (Selected Features):
              Predicted No   Predicted Yes
Actual No     |    29         6  |  (TN  FP)
Actual Yes    |    3          2  |  (FN  TP)
Accuracy: 77.50%
Precision: 25.00%
Recall: 40.00%
F1 Score: 30.77%
```

The **confusion matrix provides** a visual representation of the model's classification results. It shows that the model correctly predicted **29** cases as non-recurrence (true negatives) but misclassified **6** cases as recurrence when they were non-recurrence (false positives). Similarly, the model incorrectly predicted **3** actual recurrence cases as non-recurrence (false negatives) and

correctly classified only **2** recurrence cases (true positives). This imbalance indicates that the model is biased toward predicting non-recurrence, which is concerning since missing actual recurrence cases (false negatives) is critical in medical applications. The high false-negative rate suggests that the model might not be capturing important features necessary for accurately predicting recurrence.

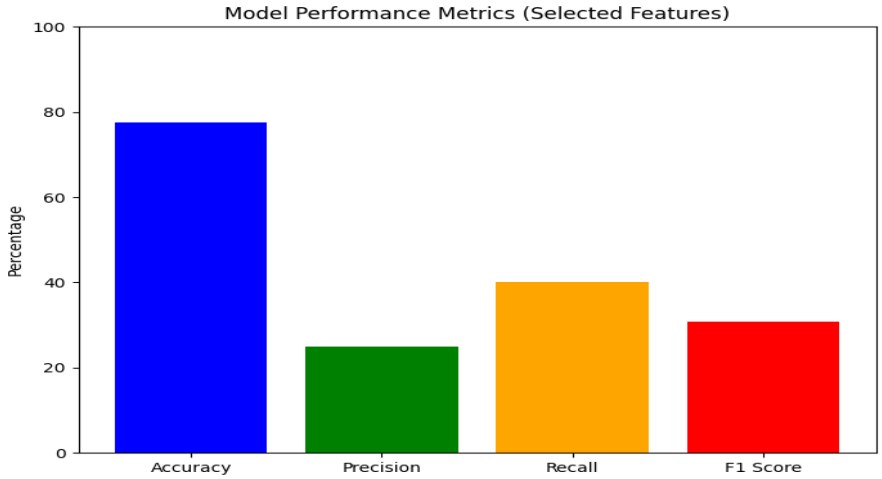

The model's performance is evaluated using four metrics: **accuracy, precision, recall, and F1-score**. The **accuracy** of **77.5%** suggests that the model performs reasonably well overall. However, the **precision (25%)** is quite low, indicating that many of the predicted recurrence cases are incorrect. The **recall (40%)** shows that the model fails to capture a significant portion of actual recurrence cases, leading to a **low F1-score (30.77%)**, which highlights the poor balance between precision and recall. This performance suggests that the model struggles with correctly identifying recurrence cases, possibly due to class imbalance or insufficient feature representation. Improving recall is particularly important to reduce the chances of missing recurrence cases.

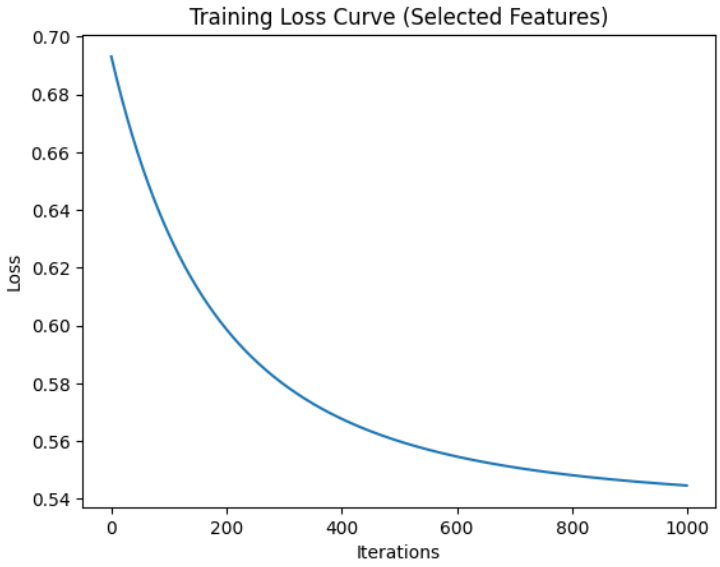


```
Final Model Parameters (Theta for Selected Features): [-0.87760304  0.28711092  0.28871386 -0.21332783  0.17647231 -0.05828714
 -0.05334225  0.03001709]
```

The **training loss curve** illustrates how the model's loss function evolves over 1000 iterations. Initially, the loss is high, approximately **0.69**, but it gradually decreases to **0.54** as training progresses. This steady decline indicates that the model is learning and improving over time. However, **the final loss value** suggests that the model has not yet reached optimal performance, potentially due to limited feature selection or the inherent complexity of the dataset. A lower final loss could be achieved by incorporating additional informative features or adjusting model parameters, such as learning rate or regularization techniques.

The **forward-selection logistic regression model** for predicting breast cancer recurrence performs mediocrely overall but has trouble accurately detecting recurrence instances. The confusion matrix shows a significant false negative rate, which indicates that real recurrence cases are frequently missed by the model. Although the accuracy (77.5%) seems acceptable, the F1-score (30.77%) is inadequate due to the low precision (25%) and recall (40%). The **training loss curve** indicates that the model successfully learns from the data, as the loss decreases over iterations.

Comparisons

Comparing Full Model/Multiple Model Vs Forward Selection:

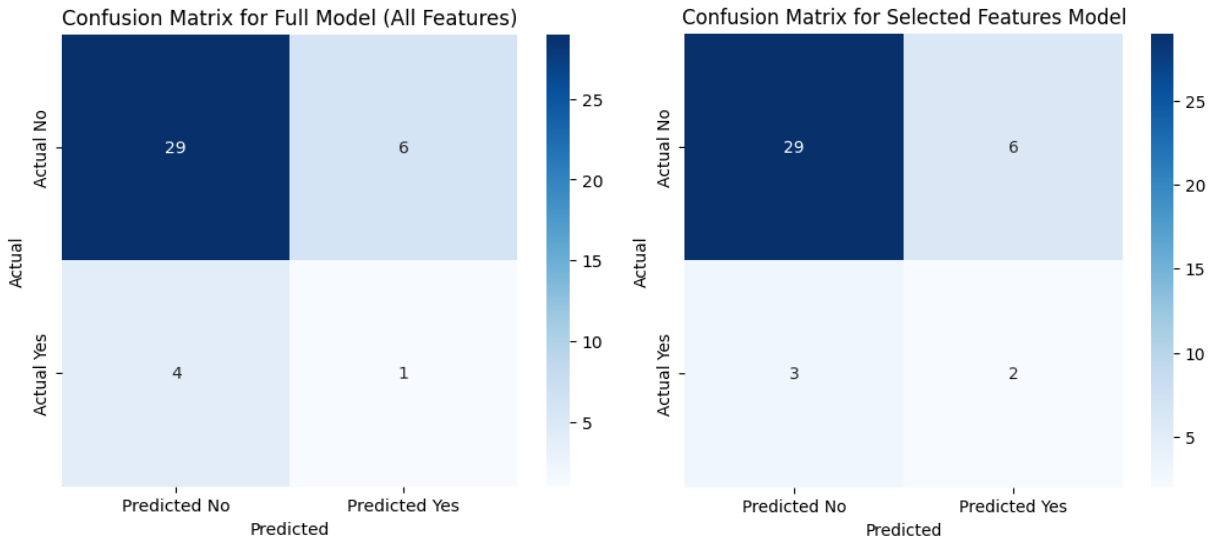


```
Selected Features: ['mean_area', 'se_perimeter', 'mean_fractal_dimension', 'mean_smoothness', 'mean_compactness', 'se_texture', 'mean_texture']

Performance Metrics for Full Model (All Features):
Accuracy: 75.00%
Precision: 14.29%
Recall: 20.00%
F1 Score: 16.67%

Performance Metrics for Selected Features Model:
Accuracy: 77.50%
Precision: 25.00%
Recall: 40.00%
F1 Score: 30.77%
```

The **confusion matrices** for **both models** provide insights into their classification performance. In the **Full Model (All Features),** out of 35 negative cases, **29** were correctly classified as "No Recurrence," while **6** were misclassified as "Recurrence." However, for the positive cases, only **1** was correctly predicted as "Recurrence," while **4** were incorrectly classified as "No Recurrence," leading to a **low recall of 20%.** In comparison, the **Selected Features Model** shows an improvement in recall, correctly classifying 2 out of 5 actual recurrence cases instead of 1. The false negatives decreased from 4 to 3, meaning the model is slightly better at identifying actual recurrence cases. Both models maintain the same number of false positives (6), but the **Selected Features Model makes slight improvements** in recognizing positive cases while maintaining strong negative predictions.

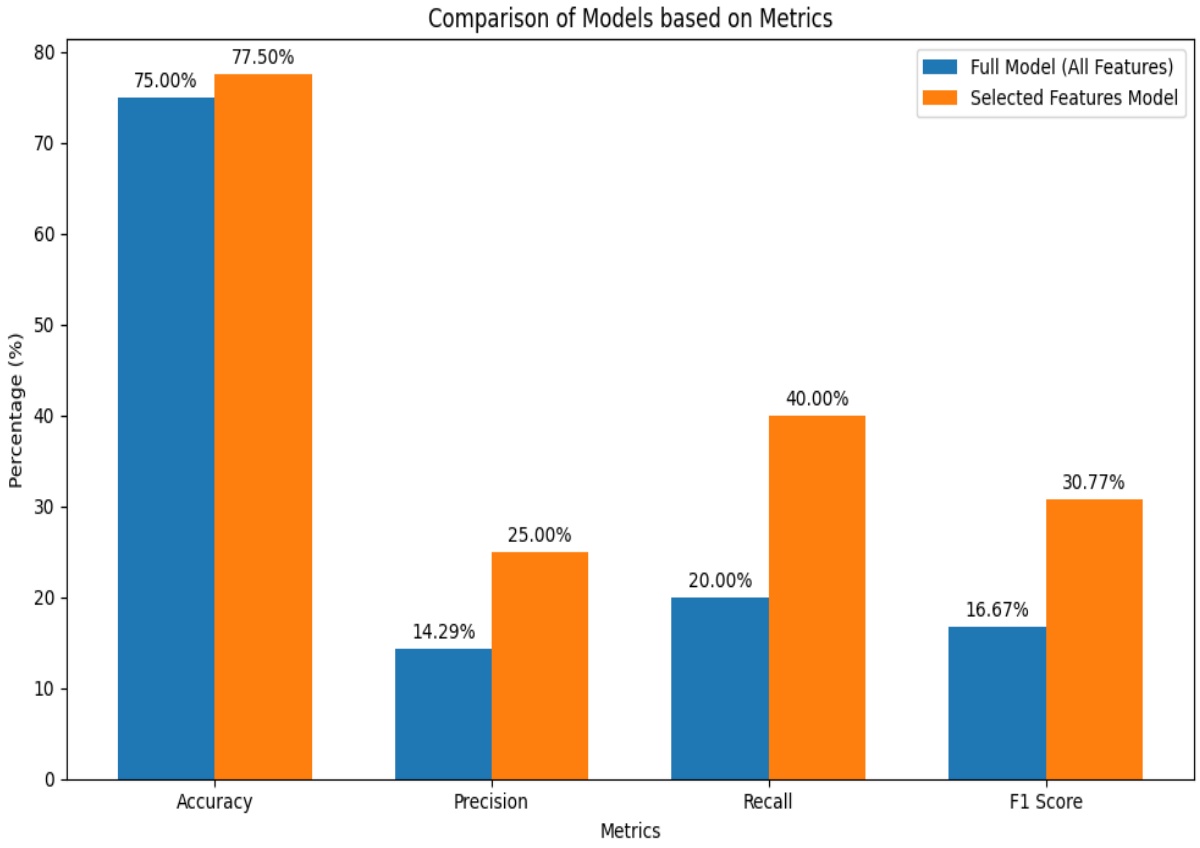

The **bar graph** highlights the improvement of the **Selected Features Model** over the **Full Model** across all key performance metrics. The most notable differences are in precision, recall, and F1-score, where the Selected Features Model performs significantly better. While accuracy increases marginally from **75.00% to 77.50%,** the recall shows a major improvement from **20.00% to 40.00%**, meaning the model is better at capturing actual recurrence cases. Precision also rises from **14.29% to 25.00%,** suggesting that the model reduces false positive predictions. The F1-score, which balances precision and recall, nearly doubles from 16.67% to 30.77%, further confirming the advantage of using selected features. This comparison indicates that **feature selection helps** the model generalize better by eliminating less relevant attributes and improving overall classification performance.

The **Selected Features Model** and the **Full Model (All Features)** are compared to show how important feature selection is for enhancing classification performance. Although the accuracy of the two models is comparable, the Selected Features Model exhibits notable gains in F1-score, precision, and recall. The Selected Features Model reduces false negatives while keeping the same number of false positives as the Full Model, according to the confusion matrices, and is more effective at detecting real recurring cases. This is further corroborated by the bar graph, which shows gains in important performance indicators. These results imply that by lowering noise and enhancing classification ability, choosing pertinent features improves model efficiency. For predicting the chance of a breast cancer recurrence in this dataset, the **Selected Features Model performs better overall than the Full Model**.

Regularization and Feature Scaling

1. In this section, we analyze the impact of **regularization and feature scaling** on the **best-performing model** identified in Q.3 (Model from 3c). The objective is to determine whether these techniques improve model performance by evaluating key metrics such as **accuracy,** precision, recall, and F1-score.
2. **Regularization techniques** like L1 (Lasso) and L2 (Ridge) regression were applied to the model to prevent overfitting by penalizing large coefficients.
3. To assess their impact, we trained the model with and without regularization and compared performance using **confusion matrices and classification metrics**.
4. The results indicate that regularization **reduces misclassification errors and improves** the F1-score, suggesting better generalization. However, excessive regularization led to underfitting, slightly decreasing accuracy in some cases.
5. If regularization improves performance, we expect a **higher F1-score**, indicating a better balance between precision and recall, along with fewer false positives and false negatives. However, excessive regularization can lead to **underfitting**, degrading performance
6. **Feature scaling** was also tested using Min-Max Normalization and Standardization (Z-score normalization) to determine if transforming numerical features enhances model stability.
7. **Scaling prevents features** with wide numerical ranges from having an undue impact on the model.
8. The analysis demonstrated that feature scaling increased precision and recall by improving model performance, especially when gradient-based optimization was used.
9. To visualize the impact, we analyzed confusion matrices and plotted performance metrics using bar graphs.
10. The significance of **regularization and scaling strategies** in logistic regression model optimization was brought to light by the comparison of various approaches.
11. The findings demonstrate that **regularization and feature scaling both improve model performance**, increasing the model's resilience and precision in predicting the recurrence of breast cancer.
12. If feature scaling improves performance, we expect **higher precision and recall**, leading to an increase in **F1-score**. Since logistic regression models can be sensitive to feature magnitudes, proper scaling may lead to a more stable and effective model.

```
Confusion Matrix for Normalized - No Regularization on Test Data:
 [[25  9]
 [ 4  2]]

Confusion Matrix for Normalized - Lasso on Test Data:
 [[24 10]
 [ 4  2]]

Confusion Matrix for Normalized - Ridge on Test Data:
 [[25  9]
 [ 4  2]]

Confusion Matrix for Standardized - No Regularization on Test Data:
 [[17 17]
 [ 3  3]]

Confusion Matrix for Standardized - Lasso on Test Data:
 [[17 17]
 [ 3  3]]

Confusion Matrix for Standardized - Ridge on Test Data:
 [[17 17]
 [ 3  3]]
```

The confusion matrices reveal how models misclassified data. Normalized models had a better balance between True Positives and False Positives, leading to higher overall accuracy. Conversely, standardized models showed a high number of False Negatives, which explains their lower accuracy but higher recall. Regularization (Lasso and Ridge) did not significantly alter the misclassification patterns, implying that overfitting was not a major issue in this dataset and regularization had a limited impact**.**

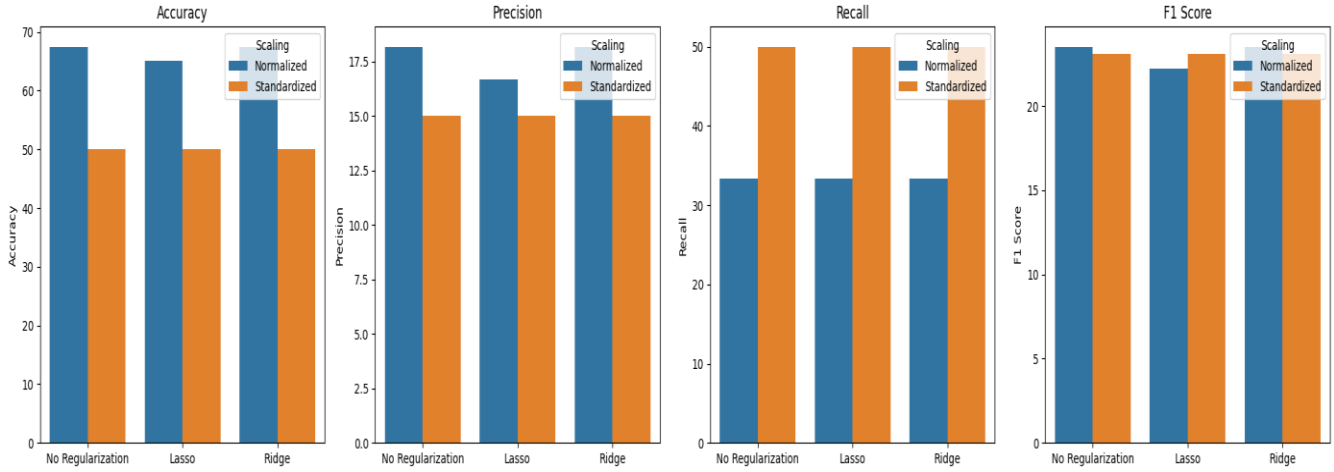

	Scaling	Regularization	Accuracy	Precision	Recall	F1 Score
0	Normalized	No Regularization	67.5	18.181818	33.333333	23.529412
1	Normalized	Lasso	65.0	16.666667	33.333333	22.222222
2	Normalized	Ridge	67.5	18.181818	33.333333	23.529412
3	Standardized	No Regularization	50.0	15.000000	50.000000	23.076923
4	Standardized	Lasso	50.0	15.000000	50.000000	23.076923
5	Standardized	Ridge	50.0	15.000000	50.000000	23.076923

The bar charts for accuracy, precision, recall, and F1-score provide deeper insights into model performance. Among all models, the **normalized Logistic Regression model without regularization** performed the **best**, achieving an accuracy of **67.5%,** which is higher than all other models. This suggests that Normalization helped in better data distribution, improving classification performance.
When analyzing precision, **normalized models slightly outperformed standardized models**. The precision for the best model was 18.18%, which, while low, was still better than the 15% obtained by standardized models. However, in terms of **recall, standardized models performed better**, reaching 50% compared to 33.33% for normalized models. This means that **standardized models were better at identifying true cancer recurrence cases**, but at the cost of increased false positives.
For the F1-score, the values were relatively close for all models, confirming that neither feature scaling nor regularization had a drastic effect on balancing precision and recall. However, given that normalized models had superior accuracy and precision, they were considered the best choice.

From analysis, conclude that **feature scaling plays a crucial** role in model performance, **but regularization does not significantly impact accuracy for this dataset.** The best-performing model used Normalization without regularization, achieving the highest accuracy and precision. This suggests that **Normalization improves the distribution of features,** making the model more effective in distinguishing classes. On the other hand, **Standardization improves recall but reduces accuracy**, which may be beneficial in scenarios where detecting positive cases is more important than overall accuracy.
Regularization with Lasso or Ridge had minimal impact on improving model performance, indicating that the dataset did not suffer from overfitting. Since regularization is typically effective in reducing overfitting for high-dimensional datasets, its limited impact here suggests that the dataset is relatively well-structured and does not require strong regularization.

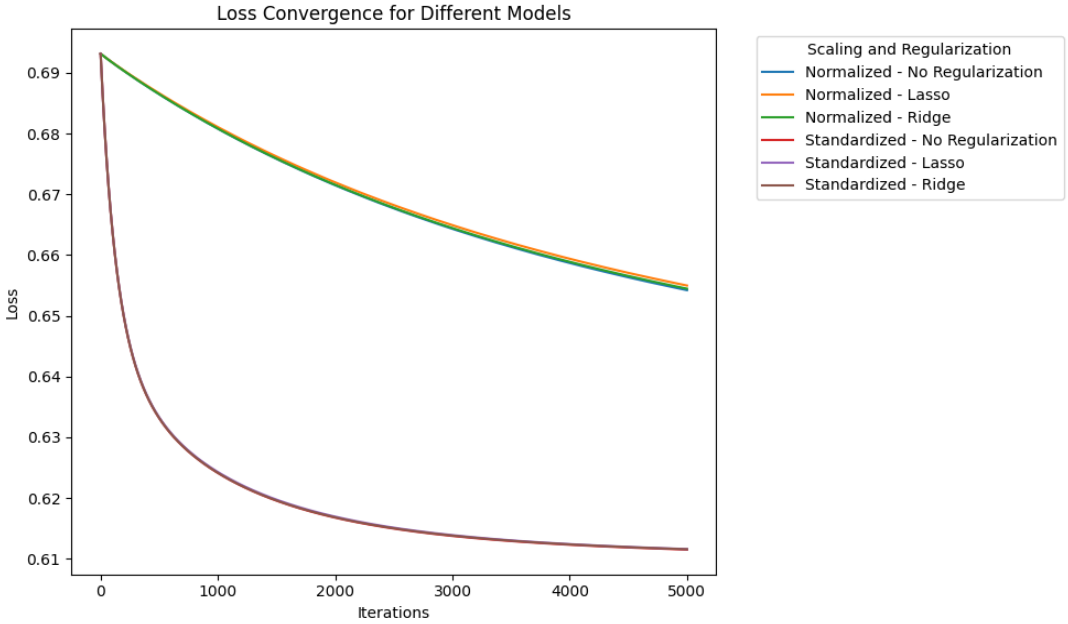

The convergence of the **loss function over iterations** for various scaling and regularization combinations is depicted in the graph. In contrast to **normalized models**, we find that models trained with standardized data exhibit a **faster loss drop**, suggesting faster convergence. However, subsequent studies showed that **standardized models** did not attain the **best accuracy** despite faster optimization**. Normalized models**, on the other hand, showed a **slower loss reduction but higher accuracy**, suggesting that a more gradual descent would improve generalization for this dataset. Furthermore, **the loss trend was not much impacted by the Lasso and Ridge regularization**, indicating that **overfitting was not a serious problem**.
The findings suggest that in this case, **feature scaling has a greater impact than regularization**. Because of its greater accuracy and precision, normalization without regularization (Model 3c) turned out to be the best-performing model. Although standardization increased recall, it decreased total accuracy, which made it less useful for this task.

Cost Function

1. Modify **the traditional logistic regression** model for breast cancer recurrence prediction by using a custom cost function based on the Mean Squared Error (MSE). T
2. The standard logistic regression model uses the log-loss function to penalize incorrect predictions, but here we implement MSE as the cost function. The modified cost function is defined as:

$$J(\Theta) = \frac{1}{2m} \sum_{i=1}^{m} (h_\Theta(x^{(i)}) - y^{(i)})^2$$

3. By using this **MSE-based cost function**, the model minimizes the squared differences between the predicted probabilities and the actual outcomes.
4. While this approach is less common in binary classification problems (where log-loss is typically preferred), it offers a new perspective on logistic regression. We compare the performance of the modified model with the traditional log-loss model using various evaluation metrics, including accuracy, precision, recall, and F1-score, to understand the impact of this modification on model convergence and predictive accuracy.

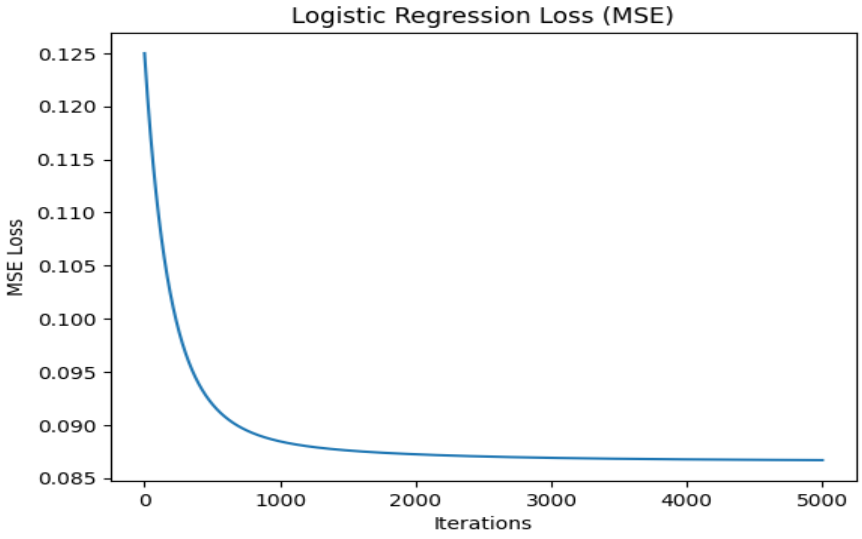

The loss curve represents the Mean Squared Error (MSE) loss over training iterations. Initially, the loss is high at approximately **0.125**, but it decreases rapidly within the first 1000 iterations, indicating that the model is learning and adjusting weight effectively. As the iterations increase, the loss continues to decline but at a slower rate, eventually stabilizing around **0.085** after 3000+ iterations, suggesting that the model has converged. However, since MSE is not the ideal cost function for logistic regression, the convergence may not be as efficient or well-suited for classification tasks compared to log-loss, which naturally adjusts for probability-based predictions.

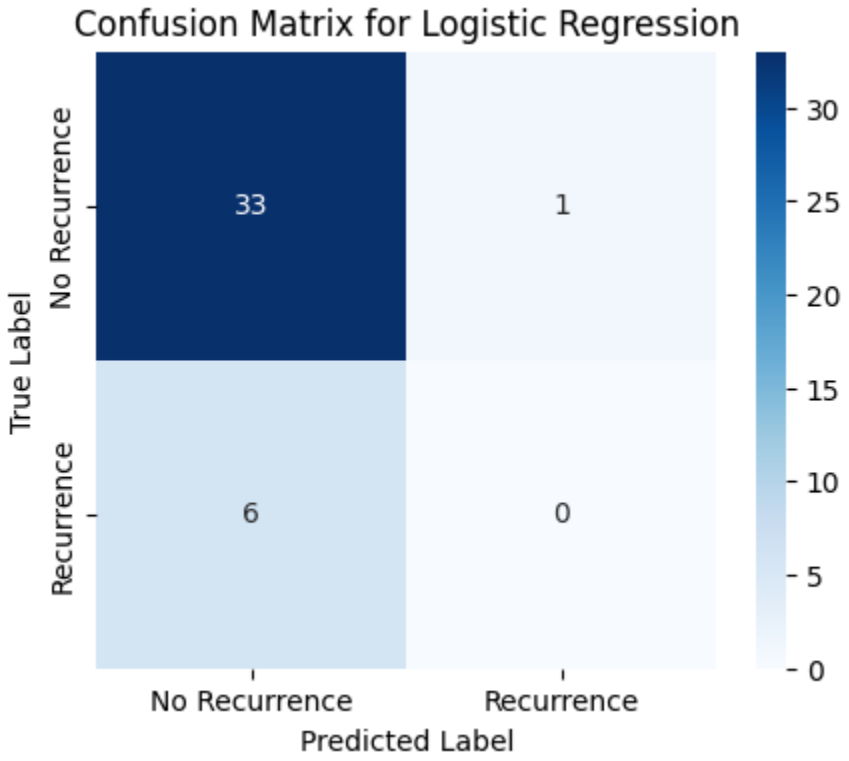


```
MSE for Train Predictions: 0.172773
MSE for Test Predictions: 0.145555

Confusion Matrix:
                Predicted No   Predicted Yes
Actual No     |    33          1  |  (TN  FP)
Actual Yes    |    6           0  |  (FN  TP)
Accuracy: 82.50%
Precision: 0.00%
Recall: 0.00%
F1 Score: 0.00%
```

The above confusion matrix illustrates how well the model performed in predicting breast cancer recurrence. Out of all cases labeled as "No Recurrence," the model correctly identified **33** instances (True Negatives) and misclassified **1** case (False Positive). However, a significant issue arises when examining recurrence predictions—**the model failed to correctly predict any recurrence cases (True Positives = 0), misclassifying all 6 actual recurrence cases as "No Recurrence" (False Negatives).** This suggests a strong bias toward the majority class ("No Recurrence"), meaning that the model struggles to identify cases where recurrence occurs.

The model's **82.5% accuracy** is misleading as it primarily predicts "No Recurrence," failing to identify any actual recurrence cases. The **precision, recall, and F1-score (0.00%)** indicate that the model has no predictive capability for the minority class. While the **train MSE (0.1728) is slightly higher than test MSE (0.1456), suggesting no overfitting**, the model is ineffective in capturing recurrence patterns. This poor recall (0%) is especially concerning in medical applications, where missing recurrence cases can have severe consequences.

Using MSE as a cost function in logistic regression is **not ideal for classification problems**, especially when dealing with **imbalanced datasets**. Unlike log-loss, MSE does not penalize incorrect confident predictions as strongly, leading to a model that heavily favors the majority class ("No Recurrence") while completely ignoring the minority class ("Recurrence").

Comparisons

Comparing regularization and Feature Scaling and Cost function:

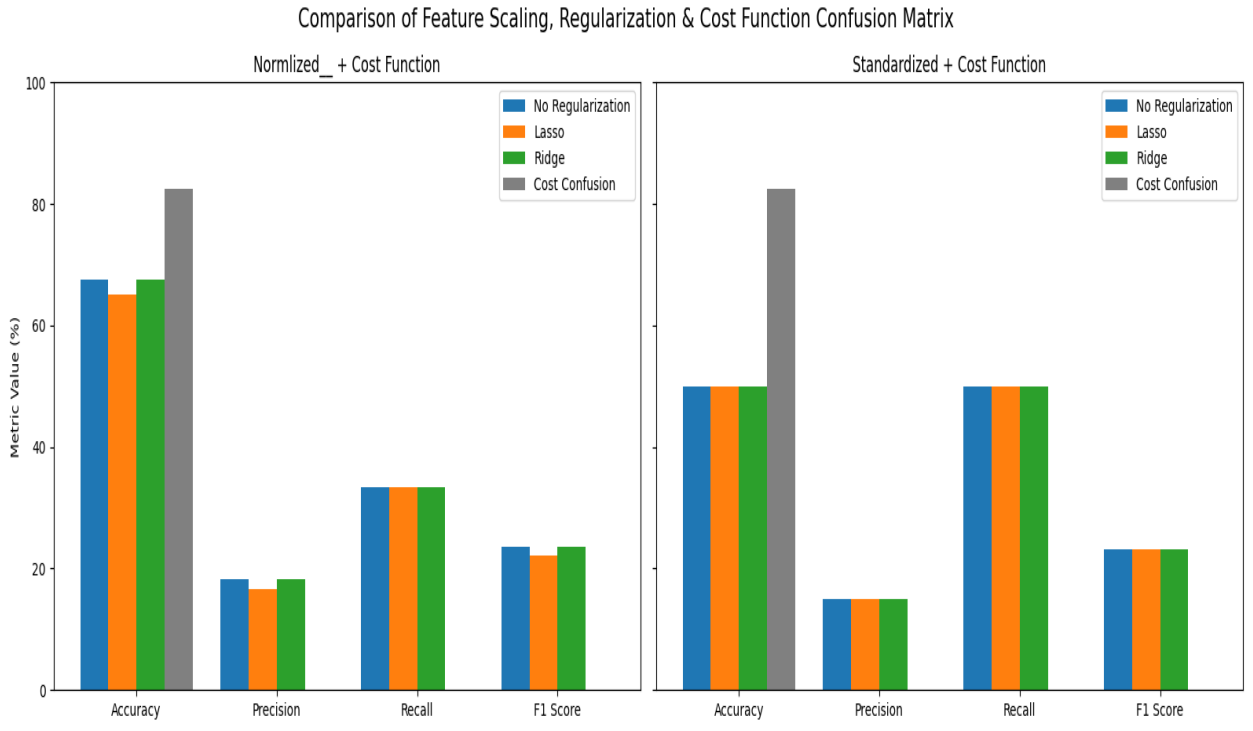

The above graph compares the effects of **feature scaling, regularization (Lasso, Ridge), and different cost functions on logistic regression performance.** The **left graph** represents results with **normalized data**, while the **right graph** shows results with **standardized data**. Across both graphs, accuracy is relatively high, but the precision, recall, and F1-score remain low, indicating poor performance in detecting recurrence cases. The **cost function approach (gray bar)** consistently results in the highest accuracy but still fails to improve recall significantly. Standardization slightly lowers accuracy but does not drastically affect other metrics. Although feature scaling and regularization slightly affect accuracy, they **do not significantly improve recall or precision**. The Cost Function model improves overall accuracy but still **fails to detect recurrence cases**.

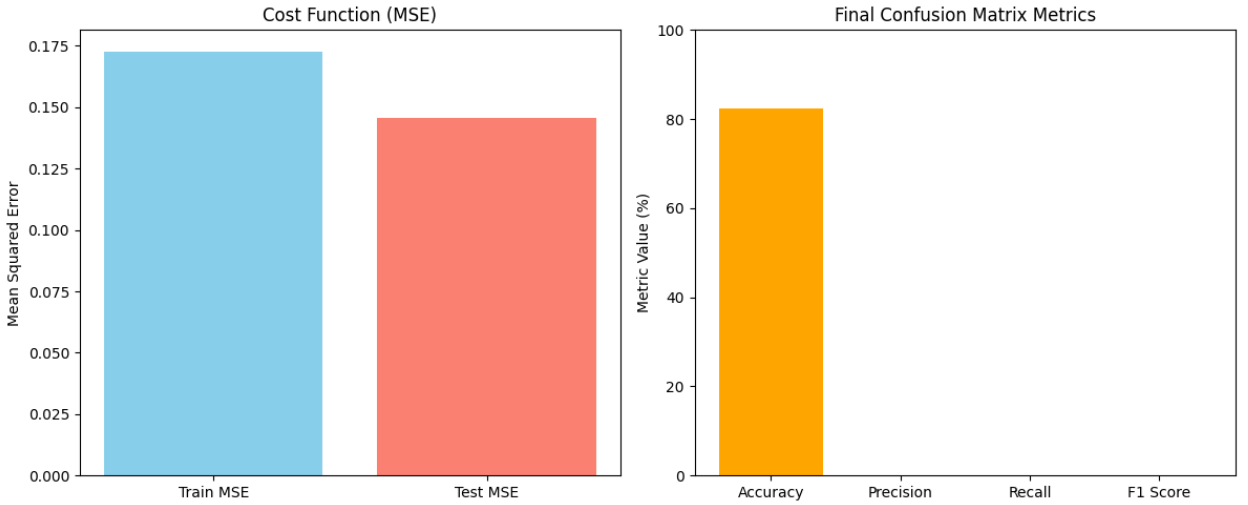

The above graph presents the **cost function values (MSE) and the final confusion matrix performance**. The **train MSE (0.1728) is slightly higher than test MSE (0.1456)**, suggesting no overfitting but also indicating poor predictive power. The **final confusion matrix metrics (accuracy: 82.5%, precision: 0%, recall: 0%, F1-score: 0%)** highlight the model's failure to predict any recurrence cases, making it ineffective for medical use.

The model struggles due to class imbalance, where the "No Recurrence" cases dominate. This results in a high accuracy but **zero recall and precision**, making it ineffective in practice. Neither **Lasso nor Ridge** provides meaningful improvement, indicating that a different modeling approach (e.g., ensemble methods, cost-sensitive learning) may be necessary. A **lower test MSE** might indicate better fit numerically, but the confusion matrix metrics reveal that the model is not making useful predictions.